

Assignment Futures 1

In this assignment, we implement futures in XINU. We have 3 modes for futures-FUTURE_EXCLUSIVE, FUTURE_SHARED and FUTURE_QUEUE.

In our implementation, futures are struct objects, defined as below:

```
typedef struct {  
    int value;  
    future_state_t state;  
    future_mode_t mode;  
    pid32 pid;  
    future_item *set_queue;  
    future_item *get_queue;  
} future_t;
```

Get_queue and Set_queue are linked lists, whose HEAD values are stored in the future. Below is the definition of the linked list items.

```
typedef struct{  
    pid32 pid;  
    struct future_item *fqnext;  
} future_item;
```

Here we are only storing the pid of the process which is waiting in the queue, and pointer for the next item in the queue.

Implementation of FUTURE_EXCLUSIVE

Implementation of this mode is done via linked lists for get_queue and set_queue. Here we can have only 1 item in get_queue, and none in set_queue.

Allocation of the Future

Allocation is done by the function future_alloc, which takes a mode, and allocates space for a new future item to be stored. It sets the state of the future as FUTURE_EMPTY, and the get_queue and set_queue values to NULL.

There are 2 functions future_get and future_set which are used to store and retrieve values to and from a future.

`future_get(future_t* f,int* address)`

Here we take the pointer of the future t as an argument, and address to which the values needs to be stored.

Here we check if the future is empty. If it is, we enqueue the process in the `get_queue`, then change the status to `FUTURE_WAITING`, and suspend it.

If it is ready, we simply retrieve the value from the future and store it in the address from the argument and change the status back to `FUTURE_EMPTY`.

Once it gets woken up(resumed) by the `future_set` function, we retrieve the value from future and store it in the address (from the argument)

If the state is already waiting, we can't do anything, since in `FUTURE_EXCLUSIVE` there can't be multiple get processes waiting for a value. So we return `SYSERR`.

`future_set(future_t* f,int value)`

Here we take the pointer of the future as and the value to put in the future as the arguments.

If the state of the future is empty, we simply store the value, change the status to `FUTURE_READY` and return.

If it is waiting, we store the value and resume the process from `get_queue` which was suspended waiting for the value we stored. There can only be 1 such process in `FUTURE_EXCLUSIVE` mode. So, if the status of the future is already in `READY` state, and future set is called, it will throw a `SYSERR`.

Implementation of `FUTURE_SHARED`

Implementation of this mode is done via linked lists for `get_queue` and `set_queue`. Here we can have multiple item in `get_queue`, and only 1 in `set_queue`.

Allocation of the Future

Allocation is done by the function `future_alloc`, which takes a mode, and allocates space for a new future item to be stored. It sets the state of the future as `FUTURE_EMPTY`, and the `get_queue` and `set_queue` values to `NULL`.

There are 2 functions `future_get` and `future_set` which are used to store and retrieve values to and from a future.

`future_get(future_t* f,int* address)`

Here we take the pointer of the future `t` as an argument, and address to which the values needs to be stored.

First we check if the future is empty. If it is, we enqueue the process in the `get_queue` and suspend it. We then change the status to `FUTURE_WAITING`.

If it is ready, we simply retrieve the value from the future and store it in the address from the argument and change the status to `FUTURE_EMPTY`.

If the state is already waiting, we traverse through the `get_queue` queue and store the process at the end of the queue. Then we suspend it.

Once it gets woken up(resumed) by the `future_set` function, we retrieve the value from future and store it in the address(from the argument)

`future_set(future_t* f,int value)`

Here we take the pointer of the future as and the value to put in the future as the arguments.

If the state of the future is empty, we simply store the value, change the status to `FUTURE_READY` and return.

If it is waiting, we store the value and resume all the waiting processes from `get_queue` by traversing the queue, which were suspended waiting for the value we just stored. There can be multiple waiting processes in the `FUTURE_SHARED` mode. After all the processes are resumed, we change back the status to `FUTURE_EMPTY`.

If the status of the future is already in `READY` state, and future set is called, it will throw a `SYSERR`.

Implementation of FUTURE_QUEUE

Implementation of this mode is done via linked lists for `get_queue` and `set_queue`. Here we can have multiple items in both `get_queue` and `set_queue`.

Allocation of the Future

Allocation is done by the function `future_alloc`, which takes a mode, and allocates space for a new future item to be stored. It sets the state of the future as `FUTURE_EMPTY`, and the `get_queue` and `set_queue` values to `NULL`.

There are 2 functions `future_get` and `future_set` which are used to store and retrieve values to and from a future.

`future_get(future_t* f,int* address)`

Here we take the pointer of the future `t` as an argument, and address to which the values needs to be stored.

First we check if the future is empty. If it is, we enqueue the process in the `get_queue` and suspend it. We then change the status to `FUTURE_WAITING`.

If it is `FUTURE_READY`, we retrieve the value from the future and store it in the address from the argument. Then we check if the `set_queue` is empty. If the `set_queue` contains processes waiting, we resume the first process from it and dequeue the process from `set_queue`. If there are no more processes waiting, we change the status to `FUTURE_EMPTY`.

If the state is already waiting, we traverse through `get_queue` and store the process at the end of the queue. Then we suspend it.

Once it gets woken up(resumed) by the `future_set` function, we retrieve the value from future and store it in the address (from the argument)

`future_set(future_t* f,int value)`

Here we take the pointer of the future as and the value to put in the future as the arguments.

If the state of the future is empty, we simply store the value and return.

If it is waiting, we store the value and resume the first process from it and dequeue the process from `get_queue`. If there are no more processes waiting, we change the status to `FUTURE_EMPTY`.

If the status is `FUTURE_READY`, we traverse through the `set_queue` queue and store the process at the end of the queue. Then we suspend it.

Free allocated memory

The function `future_free` is used to free the futures. It will check if there are any waiting processes in either of the queues, then resume the processes and free the queues. After that it frees the future object altogether.

Key points in the implementation (for all modes):

- Function `future_set` resumes processes from `get_queue`
- Function `future_get` resumes processes from `set_queue`
- `FUTURE_READY` means there can be 0 to n processes in the `set_queue` but none in the `get_queue`.
- `FUTURE_WAITING` means there can be 1 to n processes in the `get_queue` but none in the `set_queue`.
- `FUTURE_EMPTY` means there aren't any processes waiting in either queue.
- Interrupts are disabled before each implementation and restored in each function exit. (`get` and `set`)